

Indic SQA Track

Indic Spoken Question Answering — Shared Task Overview

Description

Spoken Question Answering (SQA) is the task of answering natural language questions delivered as audio. Despite significant progress for high-resource languages like English and Mandarin, Indian languages remain severely underserved — with no standardised benchmarks, evaluation protocols, or published baselines. This shared task introduces benchmark datasets for Telugu (**others -TBD**), and invites the community to build SQA systems evaluated through an LLM-as-Judge framework designed to capture the nuanced, morphologically rich nature of Indic language responses.

Target Languages

- Telugu
- **TBD**

Indic SQA Model Building

Participants are required to build a Spoken Question Answering (SQA) system for one or more of the provided Indic languages — Telugu and **others (TBD)** — using the benchmark dataset released as part of this shared task. Each system must take a spoken question (audio input) and produce a natural language answer in the source language. Participants are free to choose their own architecture, whether a pipeline approach (e.g., ASR followed by a machine reading comprehension model), an end-to-end system with speech input. System outputs will be evaluated using an **LLM-as-Judge** framework, where a large language model assesses predicted answers against gold references across multiple dimensions — factual correctness, fluency, language appropriateness, and completeness — capturing the nuanced, morphologically rich nature of Indic language responses where correct answers may be expressed in multiple valid surface forms. Participants must submit their system outputs along with a short system description paper detailing their approach, model choices, training setup, and any external resources used.

What Participants Need to Do

1	Register on the shared task portal to access the benchmark dataset.
2	Build an SQA system for one or more of the target Indic languages (Hindi, Tamil, Telugu).
3	Train and validate your system using the provided train and development splits.
4	Submit your system's predicted answers on the held-out test set via the submission portal.
5	Submit a system description paper (4–8 pages, ACL format) detailing your approach, model choices, training setup, and any external resources used.

Data

TBD

Splits provided

- Training split — released at task opening, with annotated answer spans.
- Development split — released alongside training data; gold answers included for local evaluation.
- Test split — released two weeks before the deadline; gold labels withheld and used for official scoring only.

Data is released under Creative Commons CC BY 4.0. All audio is quality-checked and manually transcribed. Transcriptions are provided as a convenience resource but participants are not required to use them.

Evaluation

Systems are evaluated using an **LLM-as-Judge** framework. A large language model scores each predicted answer against the gold reference across four dimensions:

- Factual correctness — does the answer convey the right information?
- Fluency — is the answer grammatically well-formed in the target language?
- Language appropriateness — is the response in the correct target language without unwanted code-switching?
- Completeness — does the answer fully address what was asked?

This approach is chosen over surface-level metrics like Exact Match or token F1 because Indic languages are morphologically rich and correct answers may be expressed in multiple valid surface forms. Rankings are computed per language, with an overall leaderboard averaging scores across all three languages.

Key Dates

Full dates to be confirmed. The sequence of milestones is:

- Task announcement and registration portal opens
- Training and development data released
- Test data released
- System submission deadline
- Results announced
- System description papers due
- Workshop presentation